Jan Kowalski

Python Backend & GenAI Engineer — FastAPI, LangGraph, RAG

jan.kowalski@example.com · +48 600 000 000 · Warszawa · https://www.linkedin.com/in/jan-kowalski · https://github.com/jankowalski

PROFIL

Backend engineer shipping LLM features to production: retrieval pipelines, evaluation harnesses and async APIs. Four years of Python with FastAPI, PostgreSQL and Docker on AWS. Comfortable owning a service end to end, from schema design to on-call.

KLUCZOWE TECHNOLOGIE

Python, FastAPI, LangGraph, RAG, PostgreSQL, Docker, AWS

Redis, Celery, SQLAlchemy, pytest, Kubernetes, GitHub Actions

PROJEKTY

DocuMind *Python, FastAPI, LangGraph, PostgreSQL, pgvector, Redis, Docker, AWS*

- Cut answer latency p95 from 9.1s to 2.4s with hybrid BM25 + pgvector retrieval and response streaming
- Built an offline eval harness (RAGAS + 300 labelled Q/A) that gated every prompt and model change in CI
- Served ~1.2k daily active users at 99.9% uptime over 6 months

Pricing API *Python, FastAPI, PostgreSQL, SQLAlchemy, Celery, Redis, Kubernetes*

- Sustained 1.8k req/s at p99 < 60 ms after query-plan tuning and Redis read-through caching
- Migrated a Django monolith endpoint to async FastAPI with zero downtime (shadow traffic + feature flags)

JĘZYKI

Polski (ojczysty), Angielski (C1)

Wyrażam zgodę na przetwarzanie moich danych osobowych w celu prowadzenia rekrutacji.